

On Model-Selection and Applications of Hierarchical Models

Wei Wang

Submitted in partial fulfillment of the
requirements for the degree
of Doctor of Philosophy
in the Graduate School of Arts and Sciences

COLUMBIA UNIVERSITY

2015

©2015

Wei Wang

All Rights Reserved

ABSTRACT

On Model-Selection and Applications of Hierarchical Models

Wei Wang

[illegible]

Table of Contents

List of Figures	v
List of Tables	ix
1 Introduction	1
I Cross-validation of hierarchical models	3
2 Insensitivity of Predictive Accuracy for Selecting among Multilevel Models	5
2.1 Multilevel Models and Survey Research	6
2.2 Model Assessment and Selection via Cross-Validation	7
2.2.1 Predictive Loss	7
2.2.2 Prediction Error	8
2.2.3 k -fold Cross-Validation for Estimating Predictive Loss	9
2.3 Cross-Validation of Structured Data	10
2.4 Comparing Multilevel Models for Binary Survey Outcomes	10
2.4.1 Complete Pooling, No Pooling, and Partial Pooling Models	11
2.4.2 Computation	13
2.4.3 Estimation Procedure	13
2.5 Results	14
2.5.1 Prediction Errors for a Corpus of Outcomes	14
2.5.2 How Sample Size Changes the Dynamics	18

2.5.3	Balancedness of the Hierarchical Structure	18
2.6	Discussion	21
	Bibliography	24
II	Pollings on Xbox	27
3	Hierarchical Modeling of Non-Representative Polls	29
3.1	A Brief History of Representative Sampling vs Non-representative Sampling	29
3.1.1	Xbox Data	31
3.2	Estimating voter intent with multilevel regression and poststratification . .	32
3.2.1	Multilevel regression and poststratification	32
3.2.2	National and State Voter Intent	36
3.2.3	Voter intent for demographic subgroups	37
3.3	Forecasting Election Day Outcome	38
3.3.1	Converting Voter Intent to Forecasts	38
3.3.2	National and state election day forecasts	41
3.4	Conclusion	42
	Bibliography	43
III	Causal Inference of Meta-anlalysis	57
4	Causal Inference of Meta-analysis via Gaussian Processesf	59
4.1	Meta-Analysis	59
4.1.1	Individual-Participant Meta-Analyses	60
4.2	Potential Outcomes for Meta-Analysis	61
4.2.1	Potential Outcomes	61
4.2.2	Extended Potential Outcomes	62
4.3	Meta-Analysis using Bayesian Non-parametrics	64
4.3.1	Gaussian Processes	64
4.3.2	Inference for Standard GP	65
4.3.3	Gaussian Process with Hierarchical Structure	66

4.3.4	Network Meta-Analysis	67
4.4	Models Descriptions	67
4.5	Inference	67
4.6	REAL DATA (STAR)	67
	Bibliography	69
IV	Conclusions	71
5	Conclusions	73

List of Figures

- 2.1 *Measure of fit (estimated prediction error) for all response outcomes in the 2006 Cooperative Congressional Election Survey. Outcomes are ordered by the lower bound (in-sample loss of the saturated model). The no pooling model gives a bad fit. Partial pooling does best but in most cases is almost indistinguishable from complete pooling under the cross-validation criterion.* 15

- 2.2 *Left panel: Cell proportion estimates for three models of vote intention. Each line is a state. The partial pooling model pools so much that it is indistinguishable from complete pooling. Right panel: The same estimates for the 10 most populous states. Still, partial pooling estimates are similar to complete pooling estimates.* 17

- 2.3 *Estimated prediction error of all response outcomes for augmented data sets. From top to bottom, the data sets have 2, 3, and 4 times as many data points as the original data set. The outcomes are ordered by the in-sample predictive loss. As sample size grows, complete pooling gradually gets worse and no pooling gets better.* 19

- 2.4 *Prediction error of the three models as sample size grows. The outcome under consideration is partisan vote preference in the upcoming congressional election. By this criterion, partial pooling and complete pooling perform similarly until sample size exceeds 50,000.* 20

2.5	<i>Measure of fit (prediction error) for all outcomes, ordered by in-sample training loss. The data set is simulated from real data set, and has the same sample size in total as the real data set, but keeping all demographic-geographic cells balanced. In this case, complete pooling model has much higher prediction errors than no pooling and partial pooling. Partial pooling is slightly but consistently better than no pooling. In particular, no pooling model has huge prediction error for outcomes that have smaller in-sample training loss. . .</i>	22
2.6	<i>Prediction error of the three models as sample size grows under the simulated balanced data set. The outcome under consideration is the vote for the Republican candidate in the U.S House of Representatives. Partial pooling has the lowest prediction error when sample size is under 70,000.</i>	23
3.1	A comparison of the demographic, partisan, and 2008 vote distribution in the Xbox dataset and the 2012 electorate (as measured by adjusted exit polls). The sex and age distributions, as one might expect, exhibit considerable differences.	45
3.2	Daily (unadjusted) Xbox estimates of two-party Obama support during the 45 days leading up to the 2012 presidential election, which suggest a landslide victory for Mitt Romney. The dotted blue line indicates a consensus average of traditional polls (the daily aggregated polling results from Pollster.com), the horizontal dashed line at 52% indicates the actual two-party vote share obtained by Barack Obama, and the vertical dotted lines give the dates of the three presidential debates.	46

3.3	National MRP-adjusted voter intent of two-party Obama support over the 45-day period and the associated 95% confidence bands. The horizontal dashed line indicates the actual two-party Obama vote share. The three vertical dotted lines indicate the presidential debates. Compared with the raw responses in Figure 3.2, the MRP-adjusted voter intent is much more reasonable, and voter intent in the last few days is very close to the actual outcome. For comparison, the daily aggregated polling results from Pollster.com, shown as the blue dotted line, are further away from the actual vote share than the estimates generated from the Xbox data in the last few days.	47
3.4	MRP-adjusted daily voter intent for the 12 states with the most electoral votes, and the associated 95% confidence bands. The horizontal dashed lines in each panel give the actual two-party Obama vote shares in that state. The mean and median absolute errors of the last day voter intent across the 51 Electoral College races are 2.5 and 1.8 percentage points, respectively. The state-by-state daily aggregated polling results from Pollster.com, given in the dotted blue lines, are broadly consistent with the estimates from the Xbox data.	48
3.5	Comparison of two-party Obama vote share for various demographic subgroups, as estimated from the 2012 national exit poll and from the Xbox data on the day before the election.	49
3.6	Two-party Obama support as estimated from the 2012 national exit poll and from the Xbox data on the day before the election, for various two-way interaction demographic subgroups (e.g., 65+ year-old women). The sizes of the dots are proportional to the population sizes of the corresponding subgroups. Subgroups within the same two-way interaction category (e.g., age by sex) have the same color.	50

3.7	Differences between the Xbox MRP-adjusted estimates and the exit poll estimates for the 30 largest two-dimensional demographic subgroups, ordered by the difference. Positive values indicate the Xbox estimate is larger than the corresponding exit poll estimate. Among these 30 subgroups, the median and mean absolute differences are 1.9 and 2.2 percentage points, respectively.	51
3.8	Projected Obama share of the two-party vote on election day for each of the 12 states with the most electoral votes, and associated 95% confidence bands. Compared to the MRP-adjusted voter intent in Figure 3.4, the projected two-party Obama support is more stable, and the North Carolina race switches direction after applying the calibration model. Additionally, the confidence bands become much wider and give more reasonable state-by-state probabilities of Obama victories.	52
3.9	Comparison between the probability of Obama winning the 12 largest Electoral College races based on Xbox data and on prediction market data. The prediction market data are the average of the raw Betfair and Intrade prices from winner-take-all markets. The three vertical lines represent the dates of three presidential debates. The shaded halves indicate the direction that race went.	53
3.10	Daily projections of Obama electoral votes in the 45-day period leading up to the 2012 election and associated 95% confidence bands. The solid line represents the median of the daily distribution. The horizontal dashed line represents the actual electoral votes, 332, that Obama captured in 2012 election. Three vertical dotted lines indicate the dates of three presidential debates.	54
3.11	Projected distribution of electoral votes for Obama one day before the election. The green vertical dotted line represents 269, the minimum number of electoral votes that Obama needs for a tie. The blue vertical dashed line gives 332, the actual number of electoral votes captured by Obama. The estimated likelihood of Obama winning the electoral vote is 88%.	55

List of Tables

Acknowledgments

The acknowledgments go here. The acknowledgments go here. The acknowledgments go here. The acknowledgments go here. The acknowledgments go here. The acknowledgments go here. The acknowledgments go here. The acknowledgments go here.

Dedication text

Introduction

The introduction goes here. The introduction goes here. The introduction goes here. The
introduction goes here. The introduction goes here. The introduction goes here. The
introduction goes here. The introduction goes here. The introduction goes here. The
introduction goes here. The introduction goes here. The introduction goes here. The
introduction goes here. The introduction goes here. The introduction goes here. The
introduction goes here.

Part I

Cross-validation of hierarchical models

Chapter 2

Insensitivity of Predictive Accuracy for Selecting among Multilevel Models

Models selection is an integral part of any data analysis. In an ideal world, iteratively improving and comparing model fits of different specifications should be the routine of all statistical procedures, especially when developments in methodology and computation facilitate evermore sophisticated and complex models. Often, the most important question is not that whether a more complicated model is computationally tractable, but why this model is an improvement over the older and simpler ones. Multilevel models (also known as Hierarchical Models) are an example of modern statistical models, which specifically handles data with group structure, for example, a national survey data with geographic and demographic information or an educational intervention applied to different schools and neighborhoods.

The gold standard of model comparison is out-of-sample prediction accuracy, i.e., in the hypothetical case of more observations coming in, which model gives the best prediction of new case of outcomes based on new cases of predictors. Cross-validation is a perhaps the most widely-used method for estimating out-of-sample prediction error and comparison of statistical models. By fitting the model on the training data set and then evaluating it on

the hold-out testing set, the over-optimism of using data twice is avoided. Furthermore, attempts have been made to use cross-validated objective functions for statistical inference (Craven and Wahba 1978; Seeger 2008), thus integrating out-of-sample prediction error estimation and model selection into one step.

In this chapter, I will discuss several challenges I encounter in using cross-validation predictive accuracy in evaluating and selecting among multilevel models, specifically in binary classification models. The first challenge is the lack of clear protocol for the cross-validation procedure: to truly test the model, the holdout set cannot be a simple random sample of the data but instead needs to have some multilevel structure itself, so that entire groups as well as individual observations are held out. Hierarchical cross-validation can be performed in the context of particular applications (Price, Nero, and Gelman 1996) but it is not clear how best to subsample structured data for cross-validation in a general way. The second challenge is that, in multilevel models, the observed loss function for data-level cross-validation can be so close to flat that the cross-validation estimates of prediction errors under candidate models can be swamped by random fluctuations.

I focus on the second of these concerns, demonstrating the limitations of prediction error in the context of a set of multilevel models fit to a large cross-tabulated national survey. An innovative aspect of our analysis is that we evaluate separately on 71 different survey responses, taking each in turn as the outcome in a comparison of regression models. This allows us to construct a relatively large corpus of data out of a single survey.

This chapter is a joint work with Andrew Gelman, and mostly based on (Wang and Gelman 2014).

2.1 Multilevel Models and Survey Research

There are two types of survey researchers, as identified by the classic book “Survey Errors and Survey Costs” (Groves 2004), the *describers*, who “use surveys to describe characteristics of a fixed population”, and the *modelers*, who “seek to identify causes of phenomena constantly occurring in a society”. The latter group developed models to generate less biased estimates, as a result of using more data and handling more inherent structure within

the data. Multilevel models, an example of the *modeler* approach, are effective in survey research, as partial pooling can yield accurate state-level estimates from national polls (Gelman and Hill 2007). Multilevel models have been successfully applied both to representative and nonrepresentative surveys to obtain accurate small-area estimation and prediction (Fay and Herriot 1979; Ghitza and Gelman 2013; Lax and Phillips 2009; Wang et al. 2014), and the practical application of such methods is currently being actively discussed in social science research (Buttice and Highton 2013; Lax and Phillips 2013). In the present paper, we conduct model selection procedures based on k -fold cross-validation and find that under this framework, the improvement of multilevel models over classical models is surprisingly small when measured on the scale of prediction error. Furthermore, we demonstrate that this lack of notable improvement is related to the sample size and data structure by repeating the analysis on simulated data sets that vary in terms of these two factors.

Our results illustrate that under multilevel structure, it could be tricky to use cross-validation in model selection, as the size of the data and how balanced the structure is heavily affect the relative performance of the models.

In the next section, I will present a fully Bayesian model comparison framework, a preparation for the real data analysis.

2.2 Model Assessment and Selection via Cross-Validation

2.2.1 Predictive Loss

I start with a loss function $l(\tilde{y}, a)$ corresponding to the inferential action a_M based on a model M , in face of future observations \tilde{y} . The available data, typically consisting of predictors x and outcomes y , are labeled as D . The corresponding predictive loss is then,

$$PL(p^t, M, D) = E_{p^t} l(\tilde{y}, a_M) = \int l(\tilde{y}, a_M) p^t(\tilde{y}) d\tilde{y} \quad (2.1)$$

where $p^t(\cdot)$ is the true distribution from which the future observations \tilde{y} are generated.

The predictive loss is affected by the form of the action a_M , the loss function l , and the data D . For example, a_M could be the mean of the posterior predictive distribution and l the mean square error loss. However, it is often convenient and theoretically desirable to

use the whole posterior predictive distribution as the inferential action and a logarithmic loss function. In addition, using the whole posterior predictive distribution has a Bayesian justification, as it reflects the full inferential uncertainty conditional on the model (Vehtari and Ojanen 2012). Substituting the choice of a_M and l into (2.1) yields,

$$\begin{aligned} PL(p^t, M, D) &= E_{p^t}[-\log p(\tilde{y}|D, M)] \\ &= - \int p^t(\tilde{y}) \log p(\tilde{y}|D, M) d\tilde{y} \end{aligned} \quad (2.2)$$

This quantity is central to predictive model selection. The fundamental difficulty in estimating it is that the true distribution $p^t(\cdot)$ is unknown.

Another important quantity arises when we approximate the true distribution with the empirical distribution, which gives the training loss,

$$\begin{aligned} TL(M, D) &= - \int \log p(y|D, M) d\hat{F}(y) \\ &= - \frac{1}{N} \sum_{y \in D} \log p(y|D, M). \end{aligned} \quad (2.3)$$

The training loss uses the same data for both estimation and evaluation and so in general underestimates prediction error.

2.2.2 Prediction Error

With (2.2), the model selection task is straightforward. Among the candidate models, the best model under this framework is the one that minimizes the predictive loss:

$$- \min_M \int p^t(\tilde{y}) \log p(\tilde{y}|D, M) d\tilde{y}, \quad (2.4)$$

which has a lower bound, $-\int p^t(\tilde{y}) \log p^t(\tilde{y}) d\tilde{y}$, which is the entropy of the true distribution. It is often more informative to look at the excess of the predictive loss over this lower bound, as shown in (2.5). I label this quantity as the prediction error. Conceptually, the prediction error indicates how far the posterior predictive distribution is from the oracle, and it is the Kullback-Leibler divergence between the posterior predictive distribution of the candidate model and the true generative model. As its form suggests, the prediction error

is the difference between log posterior predictive density and log true predictive density, averaged over the true predictive distribution,

$$\begin{aligned} PE(p^t, M, D) &= PL(p^t, M, D) - LB(p^t) \\ &= - \int p^t(\tilde{y}) \log p(\tilde{y}|D, M) d\tilde{y} + \int p^t(\tilde{y}) \log p^t(\tilde{y}) d\tilde{y}. \end{aligned} \quad (2.5)$$

So to estimate the prediction error, we need to estimate the two terms in (2.5).

2.2.3 k -fold Cross-Validation for Estimating Predictive Loss

In the predictive framework, the central obstacle of estimating the predictive loss (2.2) is that the future observations are not available. One thread of research attempts to estimate and correct the bias introduced by reusing the sample and thus gives rise to various information criteria, whose validity hinges on a number of assumptions and simplifications. Another thread of research is to use hold-out data for testing, thus making training and testing data independent. This leads to a variety of resampling procedures, including leave-one-out cross-validation, k -fold cross-validation, Monte Carlo cross-validation, and bootstrapping. In practice, k -fold cross-validation is popular due to its computational convenience and stability (Kale, Kumar, and Vassilvitskii 2011). Formally, the k -fold cross-validation of the predictive loss is given by

$$\begin{aligned} \widehat{PL}^{\text{CV}}(M, D) &= - \frac{1}{N} \sum_{k=1}^K \sum_{i \in \text{test}_k} \log p(y_i | D^k, M) \\ &= - \frac{1}{N} \sum_{i=1}^N \log p(y_i | D^{(\setminus i)}, M), \end{aligned} \quad (2.6)$$

where D^k represents the k^{th} training set, test_k represents the k^{th} testing set under the random partition and $D^{(\setminus i)}$ denotes the training set that excludes the i^{th} observation. Because k -fold cross-validation does not use all the data, the prediction error estimates are biased, but in the cases where there are relatively few predictors, this bias is small (Burman 1989).

The practical impediment of using cross-validation is the computational burden: with k -fold cross-validation, we need to fit the model k times. However, in many cases it is possible to perform the k steps in parallel.

The problem remains of estimating the second term in (2.5), namely the lower bound of predictive loss. In this paper, we use the in-sample training loss $TL(M_s, D)$ of the saturated model M_s as the surrogate for the lower bound. So the estimated prediction error is

$$\begin{aligned}\widehat{PE}(M, D) &= \widehat{PL}^{\text{CV}}(M, D) - TL(M_s, D) \\ &= -\frac{1}{N} \sum_{i=1}^N \log p(y_i | D^{(\setminus i)}, M) + \frac{1}{N} \sum_{y \in D} \log p(y | D, M_s).\end{aligned}\tag{2.7}$$

2.3 Cross-Validation of Structured Data

Standard cross-validation assumes that data are independent and with no distributional differences between the training and testing sets. For structured data, it is not always clear how best to perform this partition. (Burman, Chow, and Nolan 1994) discusses a modification of ordinary cross-validation procedure for stationary time series. In this paper, we focus on the cross-tabulated structure, which is the characteristic of survey data with discrete responses. In an unbalanced cross-tabulated data set, simple random sampling might result in undersampling of small cells. Thus, we adopt a stratified sampling approach to guarantee that each cell is partitioned into a training part and a testing part. Another possibility is to perform a cluster sampling and train the model on some cells and test the fitted model on others. This approach is related to transfer learning (Pan and Yang 2010). In the analysis of survey data, the focus is mostly on the existing cells rather than on hypothetical new cells, and so we only discuss cross-validation using stratified sampling on structured data.

2.4 Comparing Multilevel Models for Binary Survey Outcomes

The 2006 Cooperative Congressional Election Survey, the example data set in this paper, is a national stratified sample of size 30,000 that includes a wide variety of response outcomes, thus providing an ideal setting to evaluate cross-validation. Although various demographic predictors are available in this data set, we keep our model simple by using only two

predictors, state and income. Under this setting, the multilevel model is the preferred model over no pooling (saturated model) or complete pooling (additive model). On one hand, the saturated model will trigger overfitting. On the other hand, income and state are known to have strong interactions when predicting electoral choice (Gelman et al. 2009), so the additive model must be substantively inadequate.

2.4.1 Complete Pooling, No Pooling, and Partial Pooling Models

Bayesian multilevel modeling is a natural choice for analyzing cross-tabulated data. When the data provide many explanatory variables, and thus a potentially complex cross-tabulated structure, it is difficult to model the interactions among explanatory variables in classical models, since each single cell is getting sparser and the estimates become unstable. By borrowing strength across cells, a multilevel model (or, alternatively, some other structured model such as a Gaussian process) can produce stable estimates even for cells that have few observations and thus can be viewed as a multivariate regression or interpolation procedure..

We develop our model on a simple two-way cross-tabulation of survey data, with state and income as the two explanatory variables, having J_1 and J_2 levels respectively.¹ We assume no continuous predictors in our model. Let N be the total sample size of the survey, then the array of cell counts follows a multinomial distribution,

$$\mathbf{N} \sim \text{Multinomial}(N, \mathbf{p})$$

, where

$$\mathbf{N} = (N_{j_1 j_2})_{J_1 \times J_2},$$

$$\mathbf{p} = (p_{j_1 j_2})_{J_1 \times J_2}.$$

The population is thus divided into $J_1 \times J_2$ cells. We constrain our discussion to binary outcomes. Then for a respondent in cell (j_1, j_2) , the probability that he or she gives a

¹For the 2006 Cooperative Congressional Election Survey data set, there are 50 states ($J_1 = 50$), and 5 income levels ($J_2 = 5$), including less than \$20,000, \$20,000-\$40,000, \$40,000-\$75,000, \$75,000-\$150,000, and \$150,000+.

positive response is $\pi_{j_1 j_2}$, which is modeled using logistic regression:

$$\text{logit}(\pi_{j_1 j_2}) = \mathbf{Z}\boldsymbol{\beta},$$

in which \mathbf{Z} is the covariate vector and $\boldsymbol{\beta}$ includes the main and interaction effects. Since our goal of inference is on cell proportions $\pi_{j_1 j_2}$ rather than cell assignment probabilities $p_{j_1 j_2}$, we treat $p_{j_1 j_2}$ as fixed throughout.

Under this setup, we consider three models:

- Complete pooling of interactions:

$$\pi_{j_1 j_2} = \text{logit}^{-1}(\beta_{j_1}^{\text{state}} + \beta_{j_2}^{\text{inc}})$$

- No pooling:

$$\pi_{j_1 j_2} = \text{logit}^{-1}(\beta_{j_1}^{\text{state}} + \beta_{j_2}^{\text{inc}} + \beta_{j_1 j_2}^{\text{state*inc}})$$

- Partial pooling:

$$\pi_{j_1 j_2} = \text{logit}^{-1}(\beta_{j_1}^{\text{state}} + \beta_{j_2}^{\text{inc}} + \beta_{j_1 j_2}^{\text{state*inc}})$$

with $\beta_{j_1 j_2}^{\text{state inc}} \stackrel{i.i.d.}{\sim} N(0, \sigma^2)$ where the scale parameter σ is estimated from the data (with a separate value for each survey outcome).

Although nonparametric multilevel modeling, both in the Bayesian (Hjort 2010) and the frequentist (Ruppert, Wand, and Carroll 2003) perspectives, have been under rapid development, we adopt a linear parametric specification for the multilevel model, because linear parametric models are still the standard specification, and software that fit the routine linear parametric models are widely available and easily accessible to practitioners. In the remaining sections of this paper, we compare the prediction error of these three models under various real data and simulation settings.

We recognize that multilevel models in big-data applications can be much more complicated (Ghitza and Gelman 2013); we use a relatively simple example here to explore the basic ideas.

2.4.2 Computation

Ideally we want to do full Bayesian inference on our model, but for computational reasons we are currently using an approximate marginal posterior mode estimate provided by `blme` (Dorie 2013) in R, which is an extension of the widely-used `lme4` (Bates, Maechler, and Bolker 2013) package. The `lme4` package approximately integrates out the random effects to obtain an approximate marginal MLE of the scale parameter and the fixed effects. However, modal estimates can end up on the boundary due to sampling variability (Chung et al. 2013), which in our case makes the partial pooling model reduce to complete pooling. In `blme`, the scale parameter σ is also given a gamma prior with shape parameter 2.5 and rate parameter 0. The gamma prior is used to regularize the prior of the scale and pull the estimates of the interactions away from zero, a situation that often happens in modal estimation. We have developed an R package, `mrp` (Gelman et al. 2012), to streamline the multilevel model fitting and cross-validation procedure.

2.4.3 Estimation Procedure

For each outcome, we fit a multilevel logistic regression model, with additive, fully-interacted, and multilevel models. We use 5-fold cross-validation to estimate predictive loss (using more folds gives essentially identical results). We estimate the lower bound using the training loss of the saturated model.

Under the aforementioned setting, the cross-validation loss estimate is,

$$\begin{aligned}
\widehat{PL}^{\text{CV}}(M, D) &= -\frac{1}{N} \sum_{k=1}^K \sum_{j \in \text{test}_k} \log p(y_j | D^k, M) \\
&= -\frac{1}{N} \sum_{k=1}^K \sum_{i,j} [y_{ij}^{\text{test}_k} \log \hat{\pi}_{ij}^{D^k} + (n_{ij}^{\text{test}_k} - y_{ij}^{\text{test}_k}) \log(1 - \hat{\pi}_{ij}^{D^k})] \\
&= -\frac{1}{N} \sum_{i,j} \sum_{k=1}^K [y_{ij}^{\text{test}_k} \log \hat{\pi}_{ij}^{D^k} + (n_{ij}^{\text{test}_k} - y_{ij}^{\text{test}_k}) \log(1 - \hat{\pi}_{ij}^{D^k})] \\
&= -\frac{1}{N} \sum_{i,j} [y_{ij} \overline{\log \hat{\pi}_{ij}} + (n_{ij} - y_{ij}) \overline{\log(1 - \hat{\pi}_{ij})}] \\
&= -\sum_{i,j} \frac{n_{ij}}{N} [\pi_{ij} \overline{\log \hat{\pi}_{ij}} + (1 - \pi_{ij}) \overline{\log(1 - \hat{\pi}_{ij})}],
\end{aligned}$$

in which $n_{ij}^{\text{test}_k}$ is the number of respondents in cell (i, j) of the k -th testing set, $y_{ij}^{\text{test}_k}$ is the number of respondents who answered yes in cell (i, j) of the k -th testing set, correspondingly, n_{ij} and y_{ij} are the numbers of total respondents and respondents who answered yes in cell (i, j) , $\hat{\pi}_{ij}^{D^k}$ is the estimated π_{ij} using the k -th training data set, and $\overline{\log \hat{\pi}_{ij}}$ is the weighted average log posterior proportion from each fold, $\left(\sum_{k=1}^K y_{ij}^{\text{test}_k} \log \hat{\pi}_{ij}^{D^k}\right) / y_{ij}$, and $\overline{\log(1 - \hat{\pi}_{ij})}$ has the similar form. The cross-validation loss estimate is approximately a measure of loss under cell proportion distribution $(\exp(\overline{\log \hat{\pi}_{ij}}), \exp(\overline{\log(1 - \hat{\pi}_{ij})}))$ (here we say “approximately” because these two probabilities do not in general add up to 1). The quick calculation in section 1.2 suggests that we should expect to see only small improvements in cross-validation loss even from substantively important model improvements.

2.5 Results

2.5.1 Prediction Errors for a Corpus of Outcomes

We begin by estimating the prediction errors of all outcomes in the survey. The results are shown in Figure~2.1. The x -axis is ordered by the in-sample training loss of the saturated model $TL(M_s, D)$, which we use as a surrogate for a lower bound of predictive loss. For complete pooling and partial pooling, the prediction error stays stable across different outcomes, while the no pooling model has huge prediction error for outcomes with small lower bounds. This finding makes sense since these are the settings where overfitting is




Figure 2.1: Measure of fit (estimated prediction error) for all response outcomes in the 2006 Cooperative Congressional Election Survey. Outcomes are ordered by the lower bound (in-sample loss of the saturated model). The no pooling model gives a bad fit. Partial pooling does best but in most cases is almost indistinguishable from complete pooling under the cross-validation criterion.

most severe (saturated models achieve the lowest in-sample training error). However, the difference in prediction error between complete pooling and partial pooling seems negligible. Partial pooling is giving essentially the same result as complete pooling, at least according to cross-validation on individual survey responses.

This seems to suggest that partial pooling does not have enough information to estimate cell-to-cell variation, thus giving an overly conservative estimate. Indeed, when we plot the estimates of $\pi_{j_1 j_2}$ for one particular outcome, vote preference for in the congressional election (see the left panel of Figure~2.2), the estimates from partial pooling are almost identical to those from complete pooling. Even for populous states where, because of their large sample size, the amount of partial pooling should be small, there are no major differences between estimates from partial pooling model and estimates from complete pooling model (see the right panel of Figure~2.2). This pattern is consistent across different outcomes.

Although we believe partial pooling is intrinsically better than complete pooling, it seems that the given data are not sufficient for the partial pooling model to pick up the interaction and unpool the estimates appropriately. It is a result of the particular characteristics of this data set? There are three factors determining the structure of the data that might affect the extent of pooling of the model. First is the sample size. If we increase the sample size to a sufficiently large level, the partial pooling model will be able to partially pool the estimates to an appropriate amount. As sample size grows, the no pooling model will eventually have the same performance as partial pooling, and it might be interesting to see at what point the saturated model becomes acceptable. The second factor affecting the relative performance of the different models is the size of the interactions that are being estimated, and the third factor is the level of imbalance in the hierarchical structure. Survey data classified by demographic and geographic predictors are typically highly unbalanced due to the long tails of sizes typical in taxonomic structures (Mandelbrot 1955). For example, the 2006 CCES includes 3,637 respondents from California but only 131 from Arkansas. This unbalanced structure will affect the amount of pooling performed by a multilevel model.

In the following subsections, we conduct simulations that vary sample size and the structure of the cells to investigate how these factors affect the relative performance of the three models as captured by cross-validation.

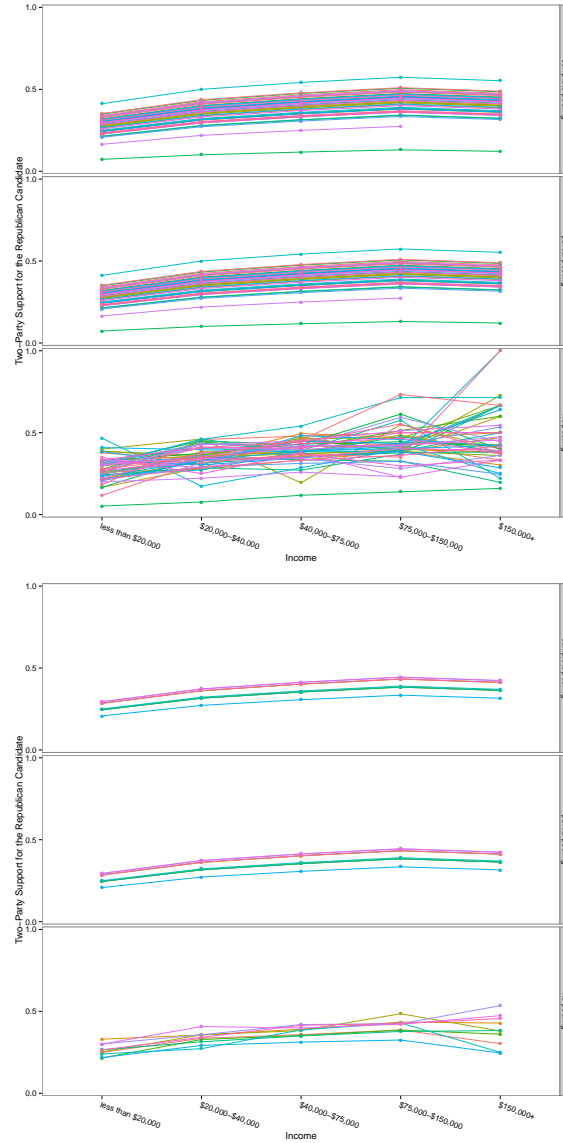


Figure 2.2: *Left panel: Cell proportion estimates for three models of vote intention. Each line is a state. The partial pooling model pools so much that it is indistinguishable from complete pooling. Right panel: The same estimates for the 10 most populous states. Still, partial pooling estimates are similar to complete pooling estimates.*

2.5.2 How Sample Size Changes the Dynamics

We artificially augment the data set by combining the data set with itself. New data sets with sample size that are 2, 3 and 4 times as large are generated. This augmentation still maintains the same level of interactions and cell structure as those of the original data. Then we estimate the prediction errors for all outcomes for the three models. Results are plotted in Figure~2.3. As we expected, as sample size grows, the prediction error of complete pooling model, which is essentially a wrong model, dominates the other two; while the prediction error of no pooling model keeps decreasing. When the sample size is 4 times as large as the original data set, no pooling model has almost the same prediction error as partial pooling model. This makes sense, since the problem of overfitting eventual goes away if we have sufficiently large sample size and fixed model structure.

These results suggest that for a fixed data structure, partial pooling decisively outperforms no pooling and complete pooling only for a certain window of sample sizes. To have a closer look at the range of the window, we look at one particular outcome, the vote preference in the upcoming election for the U.S. House of Representatives. We augment the sample size and plot the relative performance of the three models in Figure~2.4. Partial pooling model is noticeably better than complete pooling in this setup when the total sample size exceeds larger than 50,000. Other outcomes have similar patterns.

2.5.3 Balancedness of the Hierarchical Structure

One possible explanation for the steep learning curve of the partial pooling model is the highly unbalanced structure of the data. Although we have 50 states, the estimate of the covariance of the state random effects might not be reliable since some of the states have small sample sizes. To see how the balancedness of the structure affects the model, we simulate a data set based on partial pooling estimates from the original data set, but make each demographic-geographic cells of roughly the same size. The overall sample size is the same as that of the real data. Relative performance of the three models for all outcomes is plotted in Figure~2.5. The graph shows that with balanced hierarchical structure, at the same sample size and amount of interaction, partial pooling kicks in much more quickly. Thus partial pooling is consistently better than complete pooling in this scenario. As in the






Figure 2.4: Prediction error of the three models as sample size grows. The outcome under consideration is partisan vote preference in the upcoming congressional election. By this criterion, partial pooling and complete pooling perform similarly until sample size exceeds 50,000.

previous analysis, we also look at the relative performance of the three models as sample size grows. The results are plotted in Figure~2.6.

2.6 Discussion

Cross-validation is an important tool used to evaluate a wide variety of statistical methods and has been widely used in model comparison when predictive power is of concern. Some theoretical treatments have pointed out situations where cross-validation might have problems. For example, (Shao 1993) shows that, under the frequentist setting, using leave-one-out cross-validation for linear model variable selection is not consistent. However, the simplicity and transparency of cross-validation gives it a near-universal appeal. In this paper, we investigate the sensitivity of cross-validation as a model comparison instrument in a cross-tabulated multilevel survey data set.

We set up the model selection problem, considering three models for these structured data: the classical models of complete pooling and no pooling, and a Bayesian multilevel model. The multilevel model captures important interactions that are not included in the complete pooling model, while at the same time avoiding the inevitable overfitting from the no pooling model. However, the improvement of the multilevel model as given by cross-validation is surprisingly tiny, almost negligible to unsuspecting eyes. The problem is that improved fits with binary data yield minuscule improvements in log loss, in moderate sample sizes nearly indistinguishable from noise even if the improved estimates are substantively important when aggregated (for example, state-level public opinion). Simulations based on real data show that sample size and structure of the cross-tabulated cells play important roles in the relative margins of different models in cross-validation based model selection. Caution should be exercised in applying prediction error for model selection with structured data.




Figure 2.5: Measure of fit (prediction error) for all outcomes, ordered by in-sample training loss. The data set is simulated from real data set, and has the same sample size in total as the real data set, but keeping all demographic-geographic cells balanced. In this case, complete pooling model has much higher prediction errors than no pooling and partial pooling. Partial pooling is slightly but consistently better than no pooling. In particular, no pooling model has huge prediction error for outcomes that have smaller in-sample training loss.




Figure 2.6: Prediction error of the three models as sample size grows under the simulated balanced data set. The outcome under consideration is the vote for the Republican candidate in the U.S House of Representatives. Partial pooling has the lowest prediction error when sample size is under 70,000.

Bibliography

Bates, Douglas, Martin Maechler, and Ben Bolker. 2013. *Lme4: Linear mixed-effects models using s_4 classes*. <http://CRAN.R-project.org/package=lme4>.

Burman, Prabir. 1989. “A comparative study of ordinary cross-validation, v-fold cross-validation and the repeated learning-testing methods.” *Biometrika* 76(3): 503–514.

Burman, Prabir, Edmond Chow, and Deborah Nolan. 1994. “A cross-validators method for dependent data.” *Biometrika* 81(2): 351–358.

Buttice, Matthew K., and Benjamin Highton. 2013. “How does multilevel regression and poststratification perform with conventional national surveys?” *Political Analysis* 21(4): 449–467.

Chung, Yeojin et al. 2013. “A nondegenerate penalized likelihood estimator for variance parameters in multilevel models.” *Psychometrika* 78(4): 685–709.

Craven, Peter, and Grace Wahba. 1978. “Smoothing noisy data with spline functions.” *Numerische Mathematik* 31(4): 377–403.

Dorie, Vincent. 2013. *Blme: Bayesian linear mixed-effects models*. <http://CRAN.R-project.org/package=blme>.

Fay, R. E., and R. A. Herriot. 1979. “Estimates of income for small places: An application of james-stein procedures to census data.” *Journal of the American Statistical Association* 74: 269–277.

Gelman, Andrew, and Jennifer Hill. 2007. *Data analysis using regression and multi-level/Hierarchical models*. Cambridge University Press.

Gelman, Andrew et al. 2012. *Mrp: Multilevel regression and poststratification*. <http://CRAN.R-project.org/package=mrp>.

Gelman, Andrew et al. 2009. *Red state, blue state, rich state, poor state: Why americans vote the way they do, second edition*. Princeton University Press.

Ghitza, Yair, and Andrew Gelman. 2013. “Deep interactions with MRP: Election turnout and voting patterns among small electoral subgroups.” *American Journal of Political Science* 57(3): 762–776.

Groves, Robert M. 2004. *536 Survey errors and survey costs*. John Wiley & Sons.

Hjort, Nils Lid. 2010. *Bayesian nonparametrics*. Cambridge University Press.

Kale, Satyen, Ravi Kumar, and Sergei Vassilvitskii. 2011. “Cross-validation and mean-square stability.” In *Innovations in computer science*, Tsinghua University Press, p. 487–495.

Lax, Jeffrey R., and Justin H. Phillips. 2009. “How should we estimate public opinion in the states?” *American Journal of Political Science* 53(1): 107–121.

Lax, Jeffrey R., and Justin H. Phillips. 2013. “How should we estimate sub-national opinion using MRP? Preliminary findings and recommendations.” *Presented at Midwest Political Science Association*.

Mandelbrot, B. 1955. “On the language of taxonomy: An outline of a ‘thermostatistical’ theory of systems of categories with willis (natural) structure.” In *Information theory—Third london symposium*, ed. Colin Cherry., p. 135–145.

Pan, Sinno Jialin, and Qiang Yang. 2010. “A survey on transfer learning.” *IEEE Transactions on Knowledge and Data Engineering* 22(10): 1345–1359.

Price, Phillip N., Anthony V. Nero, and Andrew Gelman. 1996. “Bayesian prediction of mean indoor radon concentrations for minnesota counties.” *Health Physics* 71: 922–936.

Ruppert, David, Matt P Wand, and Raymond J Carroll. 2003. *Semiparametric regression*. Cambridge University Press.

Seeger, Matthias W. 2008. “Cross-validation optimization for large scale structured classification kernel methods.” *Journal of Machine Learning Research* 9: 1147–1178.

Shao, Jun. 1993. “Linear model selection by cross-validation.” *Journal of the American statistical Association* 88(422): 486–494.

Vehtari, Aki, and Janne Ojanen. 2012. “A survey of bayesian predictive methods for model assessment, selection and comparison.” *Statistics Surveys* 6: 142–228.

Wang, Wei, and Andrew Gelman. 2014. “Difficulty of selecting among multilevel models using predictive accuracy.” *Statistics and Its Interface* 7: 1–8.

Wang, Wei et al. 2014. “Forecasting elections with non-representative polls.” *International Journal of Forecasting*.

Part II

Pollings on Xbox

Chapter 3

Hierarchical Modeling of Non-Representative Polls

In this chapter, I will discuss an application of hierarchical modeling to non-representative survey sampling. As it is mentioned in the last chapter, there is a dichotomy in modern survey research, the camp of *describers* and the camp of *modelers*. However, at the heart of modern opinion polling, for both *describers* and *modelers*, is representative sampling, built around the goal that every individual in a particular target population (e.g., registered or likely U.S. voters) has the same probability of being sampled. Non-representative sampling has fallen out of favor among pollsters as a result of its inherent bias. I will show that, using an example of a highly-biased poll on US presidential election conducted on Xbox gaming platform, that hierarchical sampling can be used to remedy the bias and help extract useful information from non-representative polls. This is a joint work with David Rothschild, Sharad Goel, and Andrew Gelman, and has been published in (Wang et al. 2014).

3.1 A Brief History of Representative Sampling vs Non-representative Sampling

The wide-scale adoption of representative polling can largely be traced to a pivotal polling mishap in the 1936 U.S. presidential election campaign. During that campaign, the pop-

ular magazine *Literary Digest* conducted a mail-in survey that attracted over two million responses, a huge sample even by modern standards. The magazine, however, incorrectly predicted a landslide victory for Republican candidate Alf Landon over the incumbent Franklin Roosevelt. Roosevelt, in fact, decisively won the election, carrying every state except for Maine and Vermont. As pollsters and academics have since pointed out, the magazine’s pool of respondents was highly biased: it consisted mostly of auto and telephone owners as well as the magazine’s own subscribers, which underrepresented Roosevelt’s core constituencies (Squire 1988). During that same campaign, pioneering pollsters, including George Gallup, Archibald Crossley, and Elmo Roper, used considerably smaller but representative samples to predict the election outcome with reasonable (Gosnell 1937). Accordingly, non-representative or “convenience sampling” rapidly fell out of favor with polling experts. Methods used for sampling have evolved over time, from address-based, in-home interview sampling in the 1930s to random digit dialing after the growth of landlines and cellphones; nevertheless, leading polling organizations continue to put immense effort into obtaining representative samples.

Two recent trends spur the interest for non-representative polls. First, representative sampling is not nearly as representative as its name suggests, and it is becoming less so. Random digit dialing (RDD), the standard method in modern representative polling, has suffered increasingly high non-response rates, both due to the general public’s growing reluctance to answer phone surveys, and expanding technical means to screen unsolicited calls (Keeter et al. 2006). By one measure, RDD response rates have decreased from 36% in 1997 to 9% in 2012 (Kohut et al. 2012). With such low response rates, even if the initial pool of targets is representative, those who ultimately answer the phone and elect to respond are almost certainly not, calling into question the statistical benefits of such an approach. Related to dropping response rates is a corresponding increase in cost, in both time and money, as one needs to contact more and more potential respondents to find one willing to participate. The second trend driving this research is that with recent technological innovations, it is increasingly convenient and cost-effective to collect large numbers of highly non-representative samples via online surveys. What took several months for the *Literary Digest* editors to collect in 1936 can now take only a few days and

can cost just pennies per response. The challenge, of course, is to extract meaningful signal from these unconventional samples.

It is worth noting that the so-called “Big Data” is more often than not a convenient sample, with potentially huge selection bias. Without adequately addressing this issue first, any conclusion drawn from big data analysis might be misleading.

3.1.1 Xbox Data

The analysis is based on an opt-in poll continuously available on the Xbox gaming platform during the 45 days preceding the 2012 U.S. presidential election. Each day, three to five questions were posted, one of which gauged voter intention with the standard query, “If the election were held today, who would you vote for?”. Respondents were allowed to answer at most once per day. The first time they participated in an Xbox poll, respondents were additionally asked to provide basic demographic information about themselves, including their sex, race, age, education, state, party ID, political ideology, and for whom they voted in the 2008 presidential election. In total, 750,148 interviews were conducted with 345,858 unique respondents—over 30,000 of whom completed five or more polls—making this one of the largest ever election panel studies.

Despite the large sample size, the pool of Xbox respondents is far from representative of the voting population. Figure~3.1 compares the demographic composition of the Xbox participants to that of the general electorate, as estimated via the 2012 national exit poll.

¹ The most striking differences are for age and sex. As one might expect, young men domi-

¹For ease of interpretation, in Figure 3.1 states are grouped into 4 categories: (1) battleground states (Colorado, Florida, Iowa, New Hampshire, Ohio, and Virginia), the five states with the highest amounts of TV spending plus New Hampshire, which had the highest per-capita spending; (2) quasi-battleground states (Michigan, Minnesota, North Carolina, Nevada, New Mexico, Pennsylvania, and Wisconsin), which round out the states where the campaigns and their affiliates made major TV buys; (3) solid Obama states (California, Connecticut, District of Columbia, Delaware, Hawaii, Illinois, Maine, Maryland, Massachusetts, New Jersey, New York, Oregon, Rhode Island, Vermont, and Washington); and (4) solid Romney states (Alabama, Alaska, Arizona, Arkansas, Georgia, Idaho, Indiana, Kansas, Kentucky, Louisiana, Mississippi, Missouri, Montana, Nebraska, North Dakota, Oklahoma, South Carolina, South Dakota, Tennessee, Texas, Utah, West Virginia, and Wyoming).

nate the Xbox population: 18-to-29-year-olds comprise 65% of the Xbox dataset, compared to 19% in the exit poll; and men make up 93% of the Xbox sample but only 47% of the electorate. Political scientists have long observed that both age and sex are strongly correlated with voting preferences (Kaufmann and Petrocik 1999), and indeed these discrepancies are apparent in the unadjusted time-series of Xbox voter intent shown in Figure 3.2. In contrast to estimates based on traditional, representative polls (indicated by the dotted blue line in Figure 3.2), the uncorrected Xbox sample suggests a landslide victory for Mitt Romney, reminiscent of the infamous *Literary Digest* error.

3.2 Estimating voter intent with multilevel regression and poststratification

3.2.1 Multilevel regression and poststratification

To transform the raw Xbox data into accurate estimates of voter intent in the general electorate, I make use of the rich demographic information that respondents provide. In particular I *poststratify* the raw Xbox responses to mimic a representative sample of likely voters. Poststratification is a popular method for correcting for known differences between sample and target populations (Little 1993). The core idea is to partition the population into cells (e.g., based on combinations of various demographic attributes), use the sample to estimate the response variable within each cell, and finally to aggregate the cell-level estimates up to a population-level estimate by weighting each cell by its relative proportion in the population. Using y to indicate the outcome of interest, the poststratification estimate is defined by,

$$\hat{y}^{\text{PS}} = \frac{\sum_{j=1}^J N_j \hat{y}_j}{\sum_{j=1}^J N_j}$$

where \hat{y}_j is the estimate of y in cell j , and N_j is the size of the j -th cell in the population. An estimate of y can be analogously derived at any subpopulation level s (e.g., voter intent in a particular state) by

$$\hat{y}_s^{\text{PS}} = \frac{\sum_{j \in J_s} N_j \hat{y}_j}{\sum_{j \in J_s} N_j}$$

where J_s is the set of all cells that comprise s . As is readily apparent from the form of the poststratification estimator, the key is to obtain accurate cell-level estimates, as well as estimates for the cell sizes.

One of the most common ways to generate cell-level estimates is to simply average sample responses within each cell. If it is assumed that within a cell the sample is drawn at random from the larger population, this yields an unbiased estimate. However, this assumption of cell-level simple random sampling is only reasonable when the partition is sufficiently fine; on the other hand, as the partition becomes finer, the cells become sparse, and the empirical sample averages become unstable. I address these issues by instead generating cell-level estimates via a regularized regression model, namely multilevel regression.

This combined model-based poststratification strategy, known as multilevel regression and poststratification (MRP), has been used to obtain accurate small-area subgroup estimates, such as for public opinion and voter turnout in individual states and demographic subgroups [Park, Gelman, and Bafumi (2004); Lax and Phillips (2009); Ghitza and Gelman (2013)].

More formally, applying MRP in this setting comprises two steps. First a Bayesian hierarchical model is fit to obtain estimates for sparse poststratification cells; second, one averages over the cells, weighting by a measure of forecasted voter turnout, to get state and national-level estimates. Specifically, I generate the cells by considering all possible combinations of sex (2 categories), race (4 categories), age (4 categories), education (4 categories), state (51 categories), party ID (3 categories), ideology (3 categories) and 2008 vote (3 categories), which partition the data into 176,256 cells.² I fit two, nested multilevel logistic regressions to estimate candidate support in each cell. The first of the two models predicts whether a respondent supports a major-party candidate (i.e., Obama or Romney), and the second predicts support for Obama given that the respondent supports a major-

²All demographic variables are collected prior to respondents' first poll, alleviating concerns that respondents may adjust their demographic responses to be inline with their voter intention (e.g., a new Obama supporter switching his or her party ID from Republican to Democrat).

party candidate. Following the notation of (Gelman and Hill 2007), the first model is given by

$$\begin{aligned} \Pr(Y_i \in \{\text{Obama}, \text{Romney}\}) = & \\ & \text{logit}^{-1}(\alpha_0 + \alpha_1(\text{state last vote share}) \\ & + a_{j[i]}^{\text{state}} + a_{j[i]}^{\text{edu}} + a_{j[i]}^{\text{sex}} + a_{j[i]}^{\text{age}} + a_{j[i]}^{\text{race}} + a_{j[i]}^{\text{party ID}} + b_{j[i]}^{\text{ideology}} + b_{j[i]}^{\text{last vote}}) \end{aligned} \quad (3.1)$$

where α_0 is the fixed baseline intercept, and α_1 is the fixed slope for Obama's fraction of two-party vote share in the respondent's state in the last presidential election. The terms $a_{j[i]}^{\text{state}}$, $a_{j[i]}^{\text{edu}}$, $a_{j[i]}^{\text{sex}}$ and so on—which in general is denote by $a_{j[i]}^{\text{var}}$ —correspond to varying coefficients associated with each categorical variable. Here the subscript $j[i]$ indicates the cell to which the i -th respondent belongs. For example, $a_{j[i]}^{\text{age}}$ takes values from $\{a_{18-29}^{\text{age}}, a_{30-44}^{\text{age}}, a_{45-64}^{\text{age}}, a_{65+}^{\text{age}}\}$ depending on the cell membership of the i -th respondent. The varying coefficients $a_{j[i]}^{\text{var}}$ are given independent prior distributions

$$a_{j[i]}^{\text{var}} \sim N(0, \sigma_{\text{var}}^2).$$

To complete the full Bayesian specification, the variance parameters are assigned a hyperprior distribution

$$\sigma_{\text{var}}^2 \sim \text{inv-}\chi^2(\nu, \sigma_0^2),$$

with a weak prior specification for the remaining parameters, ν and σ_0 . The benefit of using a multilevel model is that estimates for relatively sparse cells can be improved through “borrowing strength” from demographically similar cells that have richer data. Similarly, the second model is defined by

$$\begin{aligned} \Pr(Y_i = \text{Obama} \mid Y_i \in \{\text{Obama}, \text{Romney}\}) = & \\ & \text{logit}^{-1}(\beta_0 + \beta_1(\text{state last vote share}) \\ & + b_{j[i]}^{\text{state}} + b_{j[i]}^{\text{edu}} + b_{j[i]}^{\text{sex}} + b_{j[i]}^{\text{age}} + b_{j[i]}^{\text{race}} + b_{j[i]}^{\text{party ID}} + b_{j[i]}^{\text{ideology}} + b_{j[i]}^{\text{last vote}}) \end{aligned} \quad (3.2)$$

and

$$b_{j[i]}^{\text{var}} \sim N(0, \eta_{\text{var}}^2),$$

$$\eta_{\text{var}}^2 \sim \text{inv-}\chi^2(\mu, \eta_0^2).$$

Jointly, Eqs. (3.1) and (3.2) define a Bayesian model that describes the data. Ideally, a fully Bayesian analysis would be performed to obtain the posterior distribution of the parameters. However, for computational convenience, I use the approximate marginal maximum likelihood estimates obtained from the `glmer()` function in the R package `lme4` (Bates, Maechler, and Bolker 2013).

Having detailed the multilevel regression step, I now turn to poststratification, where cell-level estimates are weighted by the proportion of the electorate in each cell and aggregated to the appropriate level (e.g., state or national). To compute cell weights, cross-tabulated population data is needed. One commonly used source for such data is the Current Population Survey (CPS); however, the CPS does not include some key poststratification variables, such as party identification. I thus instead use exit poll data from the 2008 presidential election. Exit polls are conducted on election day outside voting stations to record the choices of exiting voters, and they are generally used by researchers and news media to analyze the demographic breakdown of the vote (after a post-election adjustment that aligns the weighted responses to the reported state-by-state election results). In total, 101,638 respondents were surveyed in the state and national exit polls. I use the exit poll from 2008, not 2012, because this means that in theory the method as described here could have been used to generate real-time predictions during the 2012 election campaign. Admittedly, this approach puts my prediction at a disadvantage since the demographic shifts of the intervening four years cannot be captured. While combining exit poll and CPS data would arguably yield improved results, for simplicity and transparency I exclusively use the 2008 exit poll summaries for poststratification.

3.2.2 National and State Voter Intent

Figure 3.3 shows the adjusted two-party Obama support for the last 45 days of the election. The daily voter intents for two-party Obama support at the national level are illustrated in Figure 3.3. Compared with the uncorrected estimates in Figure 3.2, the MRP-adjusted estimates yield a much more reasonable timeline of Obama's standing over the course of the final weeks of the campaign. With a clear advantage at the beginning, Obama's support slipped rapidly after the first presidential debate—though never falling below 50%—and gradually recovered, building up a decisive lead in the final days.

On the day before the election, the estimate of voter intent is off by a mere 0.6 percentage points from the actual outcome (indicated by the dotted horizontal line). Voter intent in the weeks prior to the election does not directly equate to an estimate of vote share on election day—a point I return to in Section 3.3. As such, it is difficult to evaluate the accuracy of the full time-series of estimates. Nonetheless, I note that the estimates are not only intuitively reasonable, but that they are also inline with prevailing estimates based on traditional, representative polls. In particular, the estimates roughly track—and are even arguably better than—those from Pollster.com, one of the leading poll aggregators during the 2012 campaign.

National vote share receives considerable media attention, but state-level estimates are particularly relevant for many stakeholders given the role of the Electoral College in selecting the winner (Rothschild 2013). Forecasting the joint probability of victory for each candidate in state-by-state races is a challenging problem due to the interdependencies in state outcomes, and the joint electoral votes has not yet become the standard forecast. The logistical difficulties of measuring state-level vote preference, and the effort required to combine information from various sources (Lock and Gelman 2010). The MRP framework, however, provides a straightforward methodology for generating state-level results. Namely, I use the same cell-level estimates employed in the national estimate, as generated via the multilevel model in Eqs. (3.1) and (3.2), and I then poststratify to each state's demographic composition. In this manner, the Xbox responses can be used to construct estimates of voter intent over the last 45 days of the campaign for all 51 Electoral College races.

Figure 3.4 shows two-party Obama support for the 12 states with the most electoral

votes. The state timelines share similar trends (e.g., support for Obama dropping after the first debate), but also have their own idiosyncratic movements, an indication of a reasonable blend of national and state-level signals. To demonstrate the accuracy of the MRP-adjusted estimates, I plot, in dotted blue lines in Figure~3.4, the estimates generated by Pollster.com, which are broadly consistent with the state-level MRP estimates. Moreover, across the 51 Electoral College races, the mean and median absolute errors of the estimates on the day before the election are just 2.5 and 1.8 percentage points, respectively.

3.2.3 Voter intent for demographic subgroups

Apart from Electoral College races, election forecasting often focuses on candidate preference among demographic subpopulations. Such forecasts are of significant importance in modern political campaigns, which often employ targeted campaign strategies (Hillygus and Shields 2009). In the highly non-representative Xbox survey, certain subpopulations are heavily underrepresented and plausibly suffer from strong self-selection problems. This begs the question, how accurate the estimates for older women based on a platform that caters to mostly young men?

It is straightforward in MRP to estimate voter intent among any collection of demographic cells: I again use the same cell-level estimates as in the national and state settings, but poststratify to the desired target population. For example, to estimate voter intent among women, the poststratification weights are based on the relative number of women in each demographic cell. To illustrate this approach, I compute Xbox estimates of Obama support for each level of the categorical variables (e.g., males, females, Whites, Blacks, etc.) on the day before the election, and compare those with the actual voting behavior of those same groups as estimated by the 2012 national exit poll. As seen in Figure~3.5, the Xbox estimates are remarkably accurate, with a median absolute difference of 1.5 percentage points between the Xbox and the exit poll numbers.³

Not only do the Xbox data facilitate accurate estimation of voter intent across these single-dimensional demographic categories, but they also do surprisingly well at estimating two-way interactions (e.g., candidate support among 18–29 year-old Hispanics, and lib-

³Respondents' 2008 vote was not asked on the 2012 exit poll, so I exclude that comparison from Figure 3.5.

eral college graduates). Figure~3.6 shows this result, plotting the Xbox estimates against those derived from the exit polling data for each of the 149 two-dimensional demographic subgroups.⁴ Most points lie close to the diagonal, indicating that the Xbox and exit poll estimates are in agreement. Specifically, for women who are 65 and older—a group whose preferences one might a priori believe are hard to estimate from the Xbox data—the difference between Xbox and the exit poll is a mere one percentage point (49.5% and 48.5%, respectively). Across all the two-way interaction groups, the median absolute difference is just 2.4 percentage points. As indicated by the size of the points in Figure~3.6, the largest differences occur for relatively small demographic subgroups (e.g., liberal Republicans), for which both the Xbox and exit poll estimates are less reliable. For the 30 largest demographic subgroups, Figure~3.7 lists the differences between Xbox and exit poll estimates. Among these largest subgroups, the median absolute difference drops to just 1.9 percentage points.

3.3 Forecasting Election Day Outcome

3.3.1 Converting Voter Intent to Forecasts

As mentioned above, daily estimates of voter intent do not directly correspond to estimates of vote share on election day. There are two key factors for this deviation. First, opinion polls (both representative and non-representative ones) only gauge voter preference on the particular day when the poll is conducted, with the question typically phrased as, “if the election were held today.” Political scientists and pollsters have long observed that such stated preferences are prone to several biases, including the anti-incumbency bias, in which the incumbent’s polling numbers tend to be lower than the ultimate outcome (Campbell 2008), and the fading early lead bias, in which a big lead early in the campaign tends to diminish as the election gets closer (Erikson and Wlezien 2008). Moreover, voters’ attitudes

⁴State contestedness is excluded from the two-way interaction groups since the 2012 state exit polls are not yet available, and the 2012 national exit poll does not have enough data to reliably estimate state interactions; 2008 vote is also excluded, as it was not asked in the 2012 exit poll. The “other” race category was also dropped as it was not consistently defined across the Xbox and exit poll datasets.

are affected by information revealed over the course of the campaign, so preferences weeks or months before election day are at best a noisy indicator of one’s eventual vote. Second, estimates of vote share require a model of likely voters. That is, opinion polls measure preferences among a hypothetical voter pool, and are thus accurate only to the extent that this pool captures those who actually turn out to vote on election day. Both of these factors introduce significant complications in forecasting election day outcomes.

To convert daily estimates of voter intent to election day predictions—which I hereafter refer to as (???) voter intent—I compare daily voter intent in previous elections to the ultimate outcomes in those elections. Specifically, I collected historical data from three previous U.S. presidential elections, in 2000, 2004, and 2008. For each year, I obtained top-line (i.e., not individual-level) national and state estimates of voter intent from all available polls conducted in those elections.⁵ From this collection of polling data, I then constructed daily estimates of voter intent by taking a moving average of the poll numbers, in a similar manner to the major poll aggregators. Note that I rely on traditional, representative polls to reconstruct historical voter intent; in principle, however, I could have started with non-representative polls if such data were available in previous election cycles.

I next infer a mapping from voter intent to election outcomes by regressing election day vote share on the historical time-series of voter intent. The key difference between the approach in this chapter and previous related work (Erikson and Wlezien 2008; Rothschild 2009) is that I explicitly model state-level correlations, via nested national and state models and correlated error terms. Specifically, I first fit a national model given by

$$y_e^{\text{US}} = a_0 + a_1 x_{t,e}^{\text{US}} + a_2 |x_{t,e}^{\text{US}}| x_{t,e}^{\text{US}} + a_3 t x_{t,e}^{\text{US}} + \eta(t, e)$$

where y_e^{US} is the national election day vote share of the incumbent party candidate in election year e , $x_{t,e}^{\text{US}}$ is the national voter intent of the incumbent party candidate at t days before the election in year e , and $\eta \sim N(0, \sigma^2)$ is the error term. Both y_e^{US} and $x_{t,e}^{\text{US}}$ are offset by 0.5, so the values run from $-\$0.5$ to 0.5 rather than 0 to 1. The term involving the absolute value of voter intent pulls the vote share prediction toward 50%, capturing the

⁵I collected the polling data from Pollster.com and RealClearPolitics.com.

diminishing early lead effect. I do not include a main effect for time since it seems unlikely that the number of days until the election itself contributes to the final vote share directly, but rather time contributes through its interaction with the voter intent (which it is include in the model).

Similarly, the state model is given by

$$y_{s,e}^{\text{ST}} = b_0 + b_1 x_{s,t,e}^{\text{ST}} + b_2 |x_{s,t,e}^{\text{ST}}| x_{s,t,e}^{\text{ST}} + b_3 t x_{s,t,e}^{\text{ST}} + \varepsilon(s, t, e)$$

where $y_{s,e}^{\text{ST}}$ is the election day state vote share of the state's incumbent party candidate⁶ at day t , $x_{s,t,e}^{\text{ST}}$ is the state voter intent at day t , and ε is the error term. The outcome $y_{s,e}^{\text{ST}}$ is offset by the national projected vote share on that day as fit with the national calibration model, and $x_{s,t,e}^{\text{ST}}$ is offset by that day's national voter intent. Furthermore, I impose two restrictions on the magnitude and correlation structure of the error term $\varepsilon(s, t, e)$. First, since the uncertainty naturally decreases as the election gets closer (as t becomes smaller), I apply the heteroscedastic structure $\text{Var}(\varepsilon(s, t, e)) = (t + a)^2$, where a is a constant to be estimated from the data. Second, the state-specific movements within each election year are allowed to be correlated. For simplicity, and as in (Chen, Ingersoll, and Kaplan 2008), I assume these correlations are uniform (i.e., all pairwise correlations are the same), which creates one more parameter to be estimated from the data. I fit the full calibration model with the `glis()` function in the R package `nlme` (Pinheiro et al. 2012).

In summary, the procedure for generating election day forecasts proceeds in three steps:

1. Estimate the joint distribution of state and national voter intent by applying MRP to the Xbox data, as described in Section ??.
2. Fit the nested calibration model described above on historical data to obtain point estimates for the parameters, including estimates for the error terms.
3. Convert the distribution of voter intent to election day forecasts via the fitted calibration model.

⁶State incumbent parties are defined as the state-by-state winners from the previous election, which is more meaningful in this context than simply using the national incumbent.

3.3.2 National and state election day forecasts

Figure 3.8 plots the projected vote shares and pointwise 95% confidence bands over time for the 12 states with the most electoral votes. Though these time-series look quite reasonable, it is difficult to assess their accuracy as there are no ground truth estimates to compare with in the weeks prior to the election. As a starting point, I compare the state-level estimates to those generated by prediction markets, which are widely considered to be among the most accurate sources for political predictions (Rothschild 2013; Wolfers and Zitzewitz 2004). For each state, prediction markets produce daily probabilities of victory. Though Figure 3.8 plots the forecasts in terms of expected vote share, this estimation procedure in fact yields the full distribution of outcomes, and so I can likewise convert my estimates to probabilistic forecasts. Figure 3.9 shows this comparison, where the prediction market estimate is derived by averaging the two largest election markets, Betfair and Intrade. My probabilistic estimates are largely consistent with the prediction market probabilities. In fact, for races with little uncertainty (e.g., Texas and Massachusetts), the Xbox estimates do not seem to suffer from the long-shot bias common to prediction markets (Rothschild 2009), and instead yield probabilities closer to 0 or 1. For tighter races, the Xbox estimates—although still highly correlated with the prediction market probabilities—look more volatile, especially in the early part of the 45-day period. Since the ground truth is not clearly defined, it is difficult to evaluate which method—Xbox or prediction markets—yields better results. From a Bayesian perspective, if one believes the stability shown by prediction markets, this could be incorporated into the structure of the Xbox calibration model.

With the full state-level outcome distribution, I can also estimate the distribution of Electoral College votes. Figure 3.10 plots the median projected electoral votes for Obama over the last 45-days of the election, together with the 95% confidence band. In particular, on the day before the election, my model estimates Obama had an 88% chance of victory, in line with estimates based on traditional polling data. For example, Simon Jackman predicted Obama had a 91% chance of victory, using a method built from (Jackman 2005). Zooming in on the day before the election, Figure 3.11 shows the full predicted distribution of electoral votes for Obama. Compared to the actual 332 votes that Obama captured, I estimate a median of 312 votes, with the most likely outcome being 303. Though this

distribution of Electoral College outcomes seems reasonable, it does appear to have higher variance than one might expect. In particular, the extreme outcomes seem to have unrealistically high likelihood of occurring, which is likely a byproduct of the calibration model not fully capturing the state-level correlation structure. Nonetheless, given that my forecasts are based on a highly biased convenience sample of respondents, the model predictions are remarkably good.

3.4 Conclusion

Forecasts not only need to be accurate, but also relevant, timely, and cost-effective. In this chapter, I construct election forecasts satisfying all of these requirements using extremely non-representative data. Though the data were collected on a proprietary polling platform, in principle one can aggregate such non-representative samples at a fraction of the cost of conventional survey designs. Moreover, the data produce forecasts that are both relevant and timely, as they can be updated faster and more regularly than standard election polls. Thus, the key question—and one of the main contributions of this chapter—is to assess the extent to which one can generate accurate predictions from non-representative samples. Since there is limited ground truth for election forecasts, definitely establishing the accuracy of my predictions is difficult. Nevertheless, I show that the MRP-adjusted and calibrated Xbox estimates are both intuitively reasonable, and are also quite similar to those generated by more traditional means.

The greatest impact of non-representative polling will likely not be for presidential elections, but rather for smaller, local elections and specialized survey settings, where it is impractical to deploy traditional methods due to cost and time constraints. For example, non-representative polls could be used in Congressional elections, where there are currently only sparse polling data. Non-representative polls could also supplement traditional surveys (e.g., the General Social Survey) by offering preliminary results at shorter intervals. General Social Survey, which is . Finally, when there is a need to identify and track pivotal events that affect public opinion, non-representative polling offers the possibility of cost-effective continuous data collection. Standard representative polling will certainly continue to be

an invaluable tool for the foreseeable future. However, 75 years after the *Literary Digest* failure, non-representative polling (followed by appropriate post-data adjustment) is due for further exploration, for election forecasting and in social research more generally.

Bibliography

Bates, Douglas, Martin Maechler, and Ben Bolker. 2013. *Lme4: Linear mixed-effects models using s4 classes*. <http://CRAN.R-project.org/package=lme4>.

Campbell, James E. 2008. *The american campaign: US presidential campaigns and the national vote*. Texas A&M University Press.

Chen, M Keith, Jonathan E Ingersoll, and Edward H Kaplan. 2008. “Modeling a presidential prediction market.” *Management Science* 54(8): 1381–1394.

Erikson, Robert S, and Christopher Wlezien. 2008. “Are political markets really superior to polls as election predictors?” *Public Opinion Quarterly* 72(2): 190–215.

Gelman, Andrew, and Jennifer Hill. 2007. *Data analysis using regression and multi-level/hierarchical models*. Cambridge University Press.

Ghitza, Yair, and Andrew Gelman. 2013. “Deep interactions with mRP: Election turnout and voting patterns among small electoral subgroups.” *American Journal of Political Science* 57(3): 762–776.

Gosnell, Harold F. 1937. “How accurate were the polls?” *Public Opinion Quarterly* 1(1): 97–105.

Hillygus, D Sunshine, and Todd G Shields. 2009. *The persuadable voter: Wedge issues in presidential campaigns*. Princeton University Press.

Jackman, Simon. 2005. “Pooling the polls over an election campaign.” *Australian Journal of Political Science* 40(4): 499–517.

Kaufmann, Karen M, and John R Petrocik. 1999. “The changing politics of american men: Understanding the sources of the gender gap.” *American Journal of Political Science* 43(3): 864–887.

Keeter, Scott et al. 2006. “Gauging the impact of growing nonresponse on estimates from a national RDD telephone survey.” *Public Opinion Quarterly* 70(5): 759–779.

Kohut, Andrew et al. 2012. “Assessing the representativeness of public opinion surveys.” *Pew Research Center for The People & The Press* 15(May): 2012.

Lax, Jeffrey R, and Justin H Phillips. 2009. “How should we estimate public opinion in the states?” *American Journal of Political Science* 53(1): 107–121.

Little, Roderick JA. 1993. “Post-stratification: A modeler’s perspective.” *Journal of the American Statistical Association* 88(423): 1001–1012.

Lock, Kari, and Andrew Gelman. 2010. “Bayesian combination of state polls and election forecasts.” *Political Analysis* 18(3): 337–348.

Park, David K, Andrew Gelman, and Joseph Bafumi. 2004. “Bayesian multilevel estimation with poststratification: State-level estimates from national polls.” *Political Analysis* 12(4): 375–385.

Pinheiro, Jose et al. 2012. *Nlme: Linear and nonlinear mixed effects models*.

Rothschild, David. 2013. “Combining forecasts: Accurate, relevant, and timely.”

Rothschild, David. 2009. “Forecasting elections comparing prediction markets, polls, and their biases.” *Public Opinion Quarterly* 73(5): 895–916.

Squire, Peverill. 1988. “Why the 1936 literary digest poll failed.” *Public Opinion Quarterly* 52(1): 125–133.

Wang, Wei et al. 2014. “Forecasting elections with non-representative polls.” *International Journal of Forecasting*.

Wolfers, Justin, and Eric Zitzewitz. 2004. *Prediction markets*. National Bureau of Economic Research.




Figure 3.1: A comparison of the demographic, partisan, and 2008 vote distribution in the Xbox dataset and the 2012 electorate (as measured by adjusted exit polls). The sex and age distributions, as one might expect, exhibit considerable differences.




Figure 3.2: Daily (unadjusted) Xbox estimates of two-party Obama support during the 45 days leading up to the 2012 presidential election, which suggest a landslide victory for Mitt Romney. The dotted blue line indicates a consensus average of traditional polls (the daily aggregated polling results from Pollster.com), the horizontal dashed line at 52% indicates the actual two-party vote share obtained by Barack Obama, and the vertical dotted lines give the dates of the three presidential debates.




Figure 3.3: National MRP-adjusted voter intent of two-party Obama support over the 45-day period and the associated 95% confidence bands. The horizontal dashed line indicates the actual two-party Obama vote share. The three vertical dotted lines indicate the presidential debates. Compared with the raw responses in Figure 3.2, the MRP-adjusted voter intent is much more reasonable, and voter intent in the last few days is very close to the actual outcome. For comparison, the daily aggregated polling results from Pollster.com, shown as the blue dotted line, are further away from the actual vote share than the estimates generated from the Xbox data in the last few days.




Figure 3.4: MRP-adjusted daily voter intent for the 12 states with the most electoral votes, and the associated 95% confidence bands. The horizontal dashed lines in each panel give the actual two-party Obama vote shares in that state. The mean and median absolute errors of the last day voter intent across the 51 Electoral College races are 2.5 and 1.8 percentage points, respectively. The state-by-state daily aggregated polling results from Pollster.com, given in the dotted blue lines, are broadly consistent with the estimates from the Xbox data.




Figure 3.5: Comparison of two-party Obama vote share for various demographic subgroups, as estimated from the 2012 national exit poll and from the Xbox data on the day before the election.




Figure 3.6: Two-party Obama support as estimated from the 2012 national exit poll and from the Xbox data on the day before the election, for various two-way interaction demographic subgroups (e.g., 65+ year-old women). The sizes of the dots are proportional to the population sizes of the corresponding subgroups. Subgroups within the same two-way interaction category (e.g., age by sex) have the same color.




Figure 3.7: Differences between the Xbox MRP-adjusted estimates and the exit poll estimates for the 30 largest two-dimensional demographic subgroups, ordered by the difference. Positive values indicate the Xbox estimate is larger than the corresponding exit poll estimate. Among these 30 subgroups, the median and mean absolute differences are 1.9 and 2.2 percentage points, respectively.




Figure 3.8: Projected Obama share of the two-party vote on election day for each of the 12 states with the most electoral votes, and associated 95% confidence bands. Compared to the MRP-adjusted voter intent in Figure 3.4, the projected two-party Obama support is more stable, and the North Carolina race switches direction after applying the calibration model. Additionally, the confidence bands become much wider and give more reasonable state-by-state probabilities of Obama victories.




Figure 3.9: Comparison between the probability of Obama winning the 12 largest Electoral College races based on Xbox data and on prediction market data. The prediction market data are the average of the raw Betfair and Intrade prices from winner-take-all markets. The three vertical lines represent the dates of three presidential debates. The shaded halves indicate the direction that race went.




Figure 3.10: Daily projections of Obama electoral votes in the 45-day period leading up to the 2012 election and associated 95% confidence bands. The solid line represents the median of the daily distribution. The horizontal dashed line represents the actual electoral votes, 332, that Obama captured in 2012 election. Three vertical dotted lines indicate the dates of three presidential debates.




Figure 3.11: Projected distribution of electoral votes for Obama one day before the election. The green vertical dotted line represents 269, the minimum number of electoral votes that Obama needs for a tie. The blue vertical dashed line gives 332, the actual number of electoral votes captured by Obama. The estimated likelihood of Obama winning the electoral vote is 88%.

Part III

Causal Inference of Meta-analysis

Chapter 4

Causal Inference of Meta-analysis via Gaussian Processes

Meta-Analysis, the synthesis of evidence from multiple study sources, has become increasingly popular in fields such as education, psychology and public health. The major obstacle for meta-analysis is the interpretation and proper handling of study-by-study heterogeneity. Based on [sobel2015], which proposed a potential outcome framework for a formal causal approach to meta-analysis, I develop a high-dimensional Gaussian Process model that explicitly handles heterogeneities across studies, allows flexible modeling of response functions and admits fully probabilistic inference.

4.1 Meta-Analysis

Meta-analyses combine data from distinct but related studies for higher resolution inference and more nuanced understanding of the effect under investigation. Originally hailed in medical and education research, meta-analyses gain traction in wider academic disciplines as the awareness for open data is increasing across all scientific communities.

However, traditional meta-analyses are mostly conducted based on extracting and combining study-level effect summaries, since access to individual-participant level data tend to be inherently difficult to obtain. In this framework, researchers extract effect size estimates y_s and standard errors σ_s^2 , where study index $s \in \{1, \dots, S\}$ and S is the number of

studies. To handle effect size heterogeneity, a random effect model is typically used (DerSimonian and Laird 1986), in which all study effect sizes are random samples of a underlying population of effect sizes, i.e.

$$y_s = \mu_s + \sigma_s^2$$

$$\mu_s \sim \mathcal{N}(\mu_0, \tau^2)$$

Admittedly, meta-analyses based on study-level summary is still effective when the effects are homogeneous and different studies sample from similar population, they nevertheless are prone to well-known statistical fallacies, such as ecological bias, when the underlying populations and effects are heterogeneous, as it is often the case in real data.

4.1.1 Individual-Participant Meta-Analyses

Individual-Participant Data (IPD) Meta-Analysis is becoming increasingly common (Higgins et al. 2001). It has been argued that IPD data increases the power of analysis (Cooper and Patall 2009) and more robust to heterogeneous effects sizes and populations [cite]. To account for between study heterogeneity in treatment effects, the use of covariates and/or random effects models is often recommended (Aitkin 1999; Tudur Smith, Williamson, and Marson 2005). The random effects models can be seen as Bayesian Hierarchical Models (Gelman and Hill 2006), based on the justification that conditioned on appropriate set of covariates, both individual-level and study-level, the residual heterogeneities are exchangeable. There are mature softwares for fitting various types of Bayesian Hierarchical models, including Generalized Linear models and Proportional Hazard models. [cite lme4, coxme and stan]

Despite its convenient form and ease of inference, traditional IPD meta-analysis based on parametric Hierarchical models suffer from two problems. The first is the lack of formal causal framework. It is difficult to pinpoint the causal interpretation of the effect estimates from a traditional IPD hierarchical model. Consider the following example, education researchers try to determine the effect of a new intervention program, applied to different classroom and administered by different teachers. In this case, the heterogeneity

might come from either the different teachers or the different populations of schools, or both. Is the effect estimate averaging over teachers? Schools? Or both? What can we say about the effect for a new teacher, or a new school? The second problem is the inflexible form of the parametric model. Traditional parametric model requires explicit modeling assumptions from the researchers, which makes the model sensitive and facilitate potential cheery-picking. Non-parametric modeling allows flexible functional form and requires little manual tuning from the researchers.

4.2 Potential Outcomes for Meta-Analysis

(Sobel, Madigan, and Wang 2015) put meta-analysis on a concrete causal foundation by introducing a extended potential outcome framework. This section is joint work with Michael Sobel and David Madigan.

4.2.1 Potential Outcomes

Potential Outcomes Framework [] defines causal effects as comparisons of outcomes under hypothetical counter-factual treatment assignment. For example, with binary treatment $Z \in \{0, 1\}$, the causal effect of treatment Z on individual i can be defined as $y_i(1) - y_i(0)$. Typically, researchers are interested in estimating quantities such as the population average treatment effect (PATE)

$$E(Y(1) - Y(0)),$$

the population average treatment effect on the treated (PATT)

$$E(Y(1) - Y(0) \mid Z = 1)$$

The key assumption in causal inference is the ignorability assumption (or unconfoundedness assumption) (Rosenbaum and Rubin 1983), which states that given a set of observed covariates, the treatment assignment Z is independent of the potential outcomes $(Y(0), Y(1))$

$$Y(0), Y(1) \perp Z \mid X$$

In the case of randomized experiment, this assumption is trivially met without any covariates X . Under ignorability assumption,

$$\begin{aligned} & E(Y|X, z = 1) - E(Y|X, z = 0) \\ &= E(Y(1)|X, z = 1) - E(Y(0)|X, z = 0) \\ &= E(Y(1)|X) - E(Y(0)|X) \\ &= E(Y(1) - Y(0)|X) \end{aligned}$$

And thus causal effect can be identified from observations.

4.2.2 Extended Potential Outcomes

In the case of a meta-analysis, consisting S studies and Z treatment variants, the potential outcomes \mathbf{Y} for individual i can be defined as a matrix

$$\mathbf{Y}_i = \begin{pmatrix} y_i(1, 1) & y_i(1, 2) & \cdots & y_i(1, Z) \\ y_i(2, 1) & y_i(2, 2) & \cdots & y_i(2, Z) \\ \vdots & \vdots & \ddots & \vdots \\ y_i(S, 1) & y_i(S, 2) & \cdots & y_i(S, Z) \end{pmatrix}$$

With this notation, we can interpret some commonly discussed meta-analytical estimates in a causal way. For example, assuming there are only two level of treatment (0 and 1) and the causal comparison is the difference, *study-specific* treatment effect for study s is $E(y(z, s) - y(z', s))$. Note that this is different from *study-level* treatment effect θ_s in random effects models, which is $E(y(z, s) - y(z', s) \mid S = s)$. Below we will discuss conditions that will connect these two quantities.

In the context of meta-analyses, unconfoundedness can be recast as unconfoundedness within each study, i.e.,

$$Y(0, s), Y(1, s) \perp Z \mid X, S = s$$

However, this assumption is not sufficient for identifying causal effects in meta-analysis. One added layer for complexity of meta-analysis is the confounding of study selection. Consider an example of clinical trials. If some studies sample from mostly young patients while some other studies sample from mostly elderly patients, and the treatment is more effective on younger patients, then heterogeneities in treatment effects across studies would arise. Hierarchical models without adequately addressing this selection problems would result in misleading results.

However, study selection is not the only factor contributing to heterogeneities in treatment effects across studies. One lingering question is whether the same treatment z is implemented identically in all studies, or in another word, whether $Y_i(s_1, z) \stackrel{d}{=} Y_i(s_2, z)$ for all pair of $s_1, s_2 \in \{1, \dots, S\}$, where $\stackrel{d}{=}$ stands for equal in distribution. Consider an example of education intervention, in which interventions are carried out by teachers with various experience levels, then it is reasonable to question whether $Y_i(s_1, z) \stackrel{d}{=} Y_i(s_2, z)$ holds.

Two assumptions from (Sobel, Madigan, and Wang 2015) codify these two sources of heterogeneities.

A1. *Weak response consistency assumption for treatment z* : For any $z \in \{1, \dots, Z\}$ and any pair $s_1, s_2 \in \{1, \dots, S\}$,

$$Y_i(s_1, z) \stackrel{d}{=} Y_i(s_2, z)$$

A2. *Unconfounded study selection*:

$$\mathbf{Y} = \begin{pmatrix} y(1, 1) & y(1, 2) & \cdots & y(1, Z) \\ y(2, 1) & y(2, 2) & \cdots & y(2, Z) \\ \vdots & \vdots & \ddots & \vdots \\ y(S, 1) & y(S, 2) & \cdots & y(S, Z) \end{pmatrix} \perp S \mid X$$

From a frequentist point of view, these two assumptions cannot be distinguished from one other. Thus (Sobel, Madigan, and Wang 2015) suggests that meta-analysts first assess the plausibility of the two assumptions based on the characteristics of the studies, and typically assume one of these two to hold and then build models to see whether the heterogeneity

could be accounted for by the other assumption. From a Bayesian point of view, we can use a very general model, and encode regularization through appropriate prior distributions to allow for reasonable separation of these two sources of heterogeneities. This will be the topic of the following sections.

4.3 Meta-Analysis using Bayesian Non-parametrics

Traditionally, causal inference using potential outcomes focuses on two questions. Modeling of the treatment assignment process $p(z \mid x)$, also known as the propensity score, and modeling of the scientific process of how responses relate to treatment and covariates $p(y \mid z, x)$, also known as the response surface (See Rubin 2005 for more details). A myriad of methods based on the either treatment assignment mechanism (e.g., propensity score matching), or response surface modeling (e.g., regression), or combination of these two (e.g., the doubly-robust method), has been proposed for causal inference of observational data.

Recently, inspired by the advances in Bayesian non-parametric models (Hjort et al. 2010), (Hill 2011) proposed a model that focuses on accurately estimating the response surface using flexible Bayesian Additive Regression Trees, or BART (Chipman, George, and McCulloch 2010). Besides the well-known benefits of being robust to model misspecifications and being able to capture highly non-linear and interaction patterns, Bayesian non-parametric models provide natural and coherent posterior intervals to convey inferential uncertainty.

4.3.1 Gaussian Processes

Gaussian Processes (GP) have become a popular tool for nonparametric regression. A random function $f : \mathcal{X} \rightarrow \mathbb{R}$ is said to follow a GP process with kernel k if any finite-dimensional marginal of it is Gaussianly distributed, i.e.

$$f(\mathbf{x}) \sim \mathcal{N}(\boldsymbol{\mu}, \mathbf{K}_{\mathbf{x},\mathbf{x}}), \forall \mathbf{x} \in \mathbb{R}^d \text{ and } d$$

where $\mathbf{K}_{\mathbf{x},\mathbf{x}}$ is the Gram matrix of kernel k . The key component in a GP model is the

kernel k , a semi-definite function defined on $\mathcal{X} \times \mathcal{X}$ that encodes the structure. Judiciously choosing K is the most important part of fitting a GP model.

A large part of its popularity is probably due to the fact it can be interpreted as a generalization of linear regression with Gaussian errors, the predominant model for parametric regression. In fact, according to Mercer's Theorem (Williams and Rasmussen 2006), kernel k can be decomposed into

$$k(x, x') = \sum_{i=1}^{\infty} \lambda_i \phi_i(x) \phi_i^T(x')$$

where λ_i and ϕ_i are respective eigenvalues and eigenfunctions of kernel k with respect to a measure μ , i.e.,

$$\int k(x, x') \phi_i(x) d\mu(x) = \lambda_i \phi_i(x'),$$

Then GP can be considered as a basis expansion method that maps input x to an infinite dimensional space via the infinite series of functions $\{\phi_i(x)\}_{i=1}^{\infty}$.

4.3.2 Inference for Standard GP

Standard GP model for N observation pairs $(y_i, \mathbf{x}_i)_{i=1}^N$ is

$$\begin{aligned} y_i | f &\sim \mathcal{N}(f(\mathbf{x}_i), \sigma^2) \\ f &\sim GP(0, k) \end{aligned}$$

For a given kernel k , the marginal distribution of \mathbf{y} is

$$\mathbf{y} \sim \mathcal{N}(0, K_{\mathbf{x}, \mathbf{x}} + \sigma^2 I_N)$$

where $K_{\mathbf{x}, \mathbf{x}}$ is the Gram matrix of kernel k whose entries are $k(x_i, x_j)$. The predictive distribution at test points \mathbf{X}^* is

$$\begin{aligned} \mathbf{y}^* | \mathbf{X}^*, \mathbf{y}, \mathbf{X} &\sim \mathcal{N}(K_{\mathbf{X}^*, \mathbf{X}}(K_{\mathbf{X}, \mathbf{X}} + \sigma^2 I_N)^{-1} \mathbf{y}, \\ &K_{\mathbf{X}^*, \mathbf{X}^*} - K_{\mathbf{X}^*, \mathbf{X}}(K_{\mathbf{X}, \mathbf{X}} + \sigma^2 I_N)^{-1} K_{\mathbf{X}, \mathbf{X}^*}^T) \end{aligned}$$

For inference on hyperparameters, e.g., parameters governing the kernels, a standard practice is to maximize log marginal likelihood

$$\begin{aligned}\log p(\mathbf{y} \mid \mathbf{X}, \theta) &= \log \int p(\mathbf{y} \mid f, \mathbf{X}, \theta) p(f) df \\ &\propto -[\mathbf{y}^\top (K_{\mathbf{X}, \mathbf{X}}(\theta) + \sigma^2 I_N)^{-1} \mathbf{y} + \log |K_{\mathbf{X}, \mathbf{X}}(\theta) + \sigma^2 I_N|]\end{aligned}$$

and plug in the MAP (maximum a posteriori) $\hat{\theta}$ into the predictive distribution of new points \mathbf{X}^* .

4.3.3 Gaussian Process with Hierarchical Structure

Gaussian Process can be extended to handle group structure. A common approach from machine learning perspective is to pose this question as a multi-task learning problem (Bonilla, Chai, and Williams 2008; Yu, Tresp, and Schwaighofer 2005), in which the objective function’s values are vectors, or even matrices. In fact, in the case of the potential outcome framework of meta-analysis outlined above, the outcomes are study-by-treatment matrices. In this setting, the curse of dimensionalities become even more acute, as the sample size effectively mutiples by the number of outcomes under investigation. One remedy is to place restriction on the structure of kernels. For example, assuming the kernel is separable, the finite dimensional marginals of the vector-valued random function \mathbf{f} is a matrix normal distributed

$$\text{vec} \mathbf{f}(\mathbf{x}) \sim \mathcal{N}(\boldsymbol{\mu}, \mathbf{K}_{\mathbf{x}, \mathbf{x}}), \otimes \mathbf{U}_{\text{task}})$$

Assuming separability can significantly reduce the dimentionality of the problem, and properties of Kronecker product can be used to make the inference efficient (Gilboa, Saati, and Cunningham 2015). However, this approach works best for the scenario of “complete design”, meaning every combination of predictors values have been observed. This assumption is reasonable for areas such as computer experiments and robotics, where experiments can be carried out at pre-specified values of X’s; but it is quite impractical in fields such as education and public health, where researchers’ control over data collections is much limited.

4.3.4 Network Meta-Analysis

The compact form of Kronecker product assumes a block-design structure, i.e., the same set of x 's are observed for every combination of study and treatment. In reality, of course, this is not the case; causal inference, in particular, is about filling those holes, i.e., potential outcomes, in hypothetical combination of study and treatment. In fact, we only observe a small fraction of the array of matrices $\mathbf{y}_{ii} = 1^\infty$, namely $\frac{1}{ST}$ of all potential outcomes. However, this framework is general enough to handle a lot of particular problems.

Network Meta-analysis (Lumley 2002) deal with treatment pair comparisons that depend on indirect evidence. For example, if treatment A and treatment B are not assigned in any of the studies at the same time, and thus researchers have to resort to indirect comparisons, e.g., treatment C co-occurs with treatment A and treatment B in some of the studies. Since traditional analysis tend to handles one comparison at a time, violations of natural constraints are frequent, e.g., $AB+BC \neq BC$. More sophisticated models are proposed to handle this so-called “inconsistency” (Higgins et al. 2012), which in my opinion makes the models unnecessarily complex and is detrimental to intuitive understanding. The framework outlined in this thesis, however, naturally deal with network meta-analysis, since it considers all possible treatment at the same time and thus have those natural constraints built in.

4.4 Models Descriptions

4.5 Inference

4.6 REAL DATA (STAR)

The demonstrative example I use is the STAR (Student-Teacher Achievement Ratio) project, which began in 1985 to study the effect of early grades class size on student achievement in Tennessee. The study is a state-wide randomized experiment applied to over 7,000 pupils from 79 schools in 4 years. Each student was randomly assigned to one of three class types, class of 13 to 15 students, class of 22 to 25 students, and class of 22 to 25 students with a paid teaching aid. Outcomes of both standardized and curriculum-based tests were used

to assess the performance of those students in areas of math, reading and study skills. Classroom teachers were also randomly assigned to the classes they would teach. The interventions were initiated as the students entered school in kindergarten and continued through third grade, based on the common belief that early intervention has persistent effects well into later lives of the students. Due to its size and ambition, STAR is perhaps the most important education study in history. Unsurprisingly, numerous studies have been devoted to analyzing the STAR data, on both immediate effects, e.g., test scores at the end of the year of intervention [cite], and persistent effects, e.g., test scores several years after the intervention (Krueger 1997) or even earning as an adult (Chetty et al. 2010).

The public access data set is collected from Project STAR Web site at <http://www.heros-inc.org/star.htm>, with 11,598 observations on 47 variables, including information on the student, test scores and treatment assignment over the intervention years, information of the teacher, and school id et al. Due to its richness, STAR project data can be investigated in many different facets. For the sake of simplicity, we choose to look at just one outcome, standardized math test score, in one intervention year, the 1st grade. Thus we can focus on the meta-analytic part of the data, without being distracted by the longitudinal aspect of the data, which is a nuisance for our discussion.

Bibliography

Aitkin, Murray. 1999. “Meta-analysis by random effect modelling in generalized linear models.” *Statistics in Medicine* 18(17-18): 2343–2351.

Bonilla, Edwin, Kian Ming Chai, and Christopher Williams. 2008. “Multi-task gaussian process prediction.”

Chetty, Raj et al. 2010. *How does your kindergarten classroom affect your earnings? Evidence from project sTAR*. National Bureau of Economic Research.

Chipman, Hugh A, Edward I George, and Robert E McCulloch. 2010. “BART: Bayesian additive regression trees.” *The Annals of Applied Statistics*: 266–298.

Cooper, Harris, and Erika A Patall. 2009. “The relative benefits of meta-analysis conducted with individual participant data versus aggregated data.” *Psychological methods* 14(2): 165.

DerSimonian, Rebecca, and Nan Laird. 1986. “Meta-analysis in clinical trials.” *Controlled Clinical Trials* 7(3): 177–188.

Gelman, Andrew, and Jennifer Hill. 2006. *Data analysis using regression and multi-level/hierarchical models*. Cambridge University Press.

Gilboa, Elad, Yunus Saati, and John P Cunningham. 2015. “Scaling multidimensional inference for structured gaussian processes.” *Pattern Analysis and Machine Intelligence, IEEE Transactions on* 37(2): 424–436.

Higgins, JPT et al. 2012. “Consistency and inconsistency in network meta-analysis: Concepts and models for multi-arm studies.” *Research Synthesis Methods* 3(2): 98–110.

Higgins, Julian et al. 2001. “Meta-analysis of continuous outcome data from individual patients.” *Statistics in Medicine* 20(15): 2219–2241.

Hill, Jennifer L. 2011. “Bayesian nonparametric modeling for causal inference.” *Journal of Computational and Graphical Statistics* 20(1).

Hjort, Nils Lid et al. 2010. “Bayesian nonparametrics.” *AMC* 10: 12.

Krueger, Alan B. 1997. *Experimental estimates of education production functions*. National Bureau of Economic Research.

Lumley, Thomas. 2002. “Network meta-analysis for indirect treatment comparisons.” *Statistics in Medicine* 21(16): 2313–2324.

Rosenbaum, Paul R, and Donald B Rubin. 1983. “The central role of the propensity score in observational studies for causal effects.” *Biometrika* 70(1): 41–55.

Rubin, Donald B. 2005. “Causal inference using potential outcomes.” *Journal of the American Statistical Association* 100(469).

Sobel, Michael E, David B Madigan, and Wei Wang. 2015. “Meta-analysis: A causal framework, with application to randomized studies of vioxx.” *Technical Report, Department of Statistics, Columbia University*.

Tudur Smith, Catrin, Paula R Williamson, and Anthony G Marson. 2005. “Investigating heterogeneity in an individual patient data meta-analysis of time to event outcomes.” *Statistics in medicine* 24(9): 1307–1319.

Williams, CKI, and CE Rasmussen. 2006. *Gaussian processes for machine learning*. Cambridge: MIT Press.

Yu, Kai, Volker Tresp, and Anton Schwaighofer. 2005. “Learning gaussian processes from multiple tasks.” In *Proceedings of the 22nd international conference on machine learning*, ACM, p. 1012–1019.

Part IV

Conclusions

Conclusions

[illegible]